

Project: Investigate a Dataset (The Movie Database - TMDb)

Table of Contents

- [Introduction](#)
- [Questions I am going to ask and find](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Analysis Shortcoming and Data limitations](#)
- [Conclusions](#)

Introduction

I have selected TMDb Movie dataset which contains over 10,000 movies including user rating, authors, budget and revenue. In the following I will describe my understanding about columns of the table.

- `id` - Unique id of each movie
- `imdb_id` - Unique ID to access the movie from IMDB (<https://www.imdb.com/title/tt0369610/>)
- `popularity` - Popularity range from 0 to 33.
- `budget` - Estimated budget in dollars (precise to extracted date)
- `revenue` - Revenue in dollars (precise to extracted date)
- `original_title` - Title of the movie
- `cast` - Top actors/actresses/cast
- `homepage` - web home page address
- `director` - director name
- `tagline` - short text like search keyword
- `keywords` - movie keywords, need to split by |
- `overview` - Short description/storyline
- `runtime` - duration of the movie
- `genres` - genres need to split by |
- `production_companies` - productioners, need to split by |
- `release_date` - Released date
- `vote_count` - integer
- `vote_average` - average vote
- `budget_adj`, `revenue_adj` - The final two columns ending with “_adj” show the budget and revenue of the associated movie in terms of 2010 dollars

I would like to investigate some patterns associated with `genres`, `revenues`, `budget` and `profit`. Timeline such as `release_year` grouping is mostly important to do this task. After reviewing the dataset, we can clearly see `NaN/NULL` values exist. I am dropping related `NULL` rows. Therefore, I decided to work with separate dataset for each Research questions to preserve original one(Maybe, I will need some of columns to investigate). So, I am transforming main

dataframe to sub-dataframes. For example, `df_q1` (dataframe for research question 1) will be created from main `df` (dataframe), `df_q2` - dataframe for research question 2 and so on... Overall, you can consider `df_qN` s are cleaned dataset.

Questions I am going to ask and find.

- [Which genres are most popular from year to year?](#)
- [What kinds of properties are associated with movies that have high revenues?](#)
- [Which movies made the most profit, yearly?](#)

```
In [1]: import pandas as pd # CSV reader pandas library
import matplotlib.pyplot as plt # Matplotlib for styling
import seaborn as sns # Seaborn data visualization library
import matplotlib.pyplot as plt # Matplotlib for styling
import matplotlib.ticker as ticker #tick locators and formatters
import numpy as np # Numpy library
#to draw the graphs inline
%matplotlib inline
pd.options.mode.chained_assignment = None
```

Data Wrangling

General Properties

In the following several lines, I am reviewing the given dataset by printing several rows, all columns, describe, info, hist() and null_values checking for the reason of determining properties and relationships.

```
In [2]: ### Importing city_list.csv file into city_lists file object
df = pd.read_csv('./dataset/tmdb-movies.csv')
df.head(3)
```

Out[2]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://www.the

3 rows × 21 columns

```
In [3]: df.tail(3)
```

Out[3]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	homepage
10863	39768	tt0060161	0.065141	0	0	Beregis Avtomobilya	Innokentiy Smoktunovskiy Oleg Efremov Georgi Z...	NaN
10864	21449	tt0061177	0.064317	0	0	What's Up, Tiger Lily?	Tatsuya Mihashi Akiko Wakabayashi Mie Hama Joh...	NaN
10865	22293	tt0060666	0.035919	19000	0	Manos: The Hands of Fate	Harold P. Warren Tom Neyman John Reynolds Dian...	NaN

3 rows × 21 columns

```
In [4]: # printing columns
print('Columns : ', df.columns)

# To print the size of the table
print('Size of the table: ', df.shape)
```

```
Columns : Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',
                'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',
                'runtime', 'genres', 'production_companies', 'release_date',
                'vote_count', 'vote_average', 'release_year', 'budget_adj',
                'revenue_adj'],
              dtype='object')
Size of the table: (10866, 21)
```

Some of my reviews for necessary columns

- It seems from below that `id`, `budget_adj` and `revenue_adj` don't necessary for my Research Questions.
- It seems from below that more than 50% of the `budget` and `revenue` equals to zero.
- It seems from below that `popularity` value ranges from 0 to 33
- It seems from below that there exists `runtime` of the movies equal to zero

```
In [5]: df.describe()
```

Out[5]:

	id	popularity	budget	revenue	runtime	vote_count	vote_
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	!
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	(
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	.
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	!
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	(
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	(
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	!

It seems from below that `homepage`, `tagline` and `keywords` columns doesn't seem useful for further analysis. So, I will not consider them in my sub-dataframes. Worth to note that, length of the column `genres` is not equal to 10866. So, it means I have to consider NULL values and drop them.

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    10866 non-null  int64
 1   imdb_id               10856 non-null  object
 2   popularity            10866 non-null  float64
 3   budget               10866 non-null  int64
 4   revenue              10866 non-null  int64
 5   original_title       10866 non-null  object
 6   cast                 10790 non-null  object
 7   homepage             2936 non-null   object
 8   director             10822 non-null  object
 9   tagline              8042 non-null   object
10  keywords             9373 non-null   object
11  overview             10862 non-null  object
12  runtime              10866 non-null  int64
13  genres               10843 non-null  object
14  production_companies  9836 non-null   object
15  release_date         10866 non-null  object
16  vote_count           10866 non-null  int64
17  vote_average         10866 non-null  float64
18  release_year         10866 non-null  int64
19  budget_adj           10866 non-null  float64
20  revenue_adj          10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
In [7]: # to see wich columns have, how many NULL values
df.isna().sum()
```

```
Out[7]: id                0
imdb_id                10
popularity             0
budget                0
revenue               0
original_title         0
cast                  76
homepage              7930
director              44
tagline              2824
keywords             1493
overview              4
runtime               0
genres               23
production_companies  1030
release_date          0
vote_count            0
vote_average          0
release_year          0
budget_adj            0
revenue_adj           0
dtype: int64
```

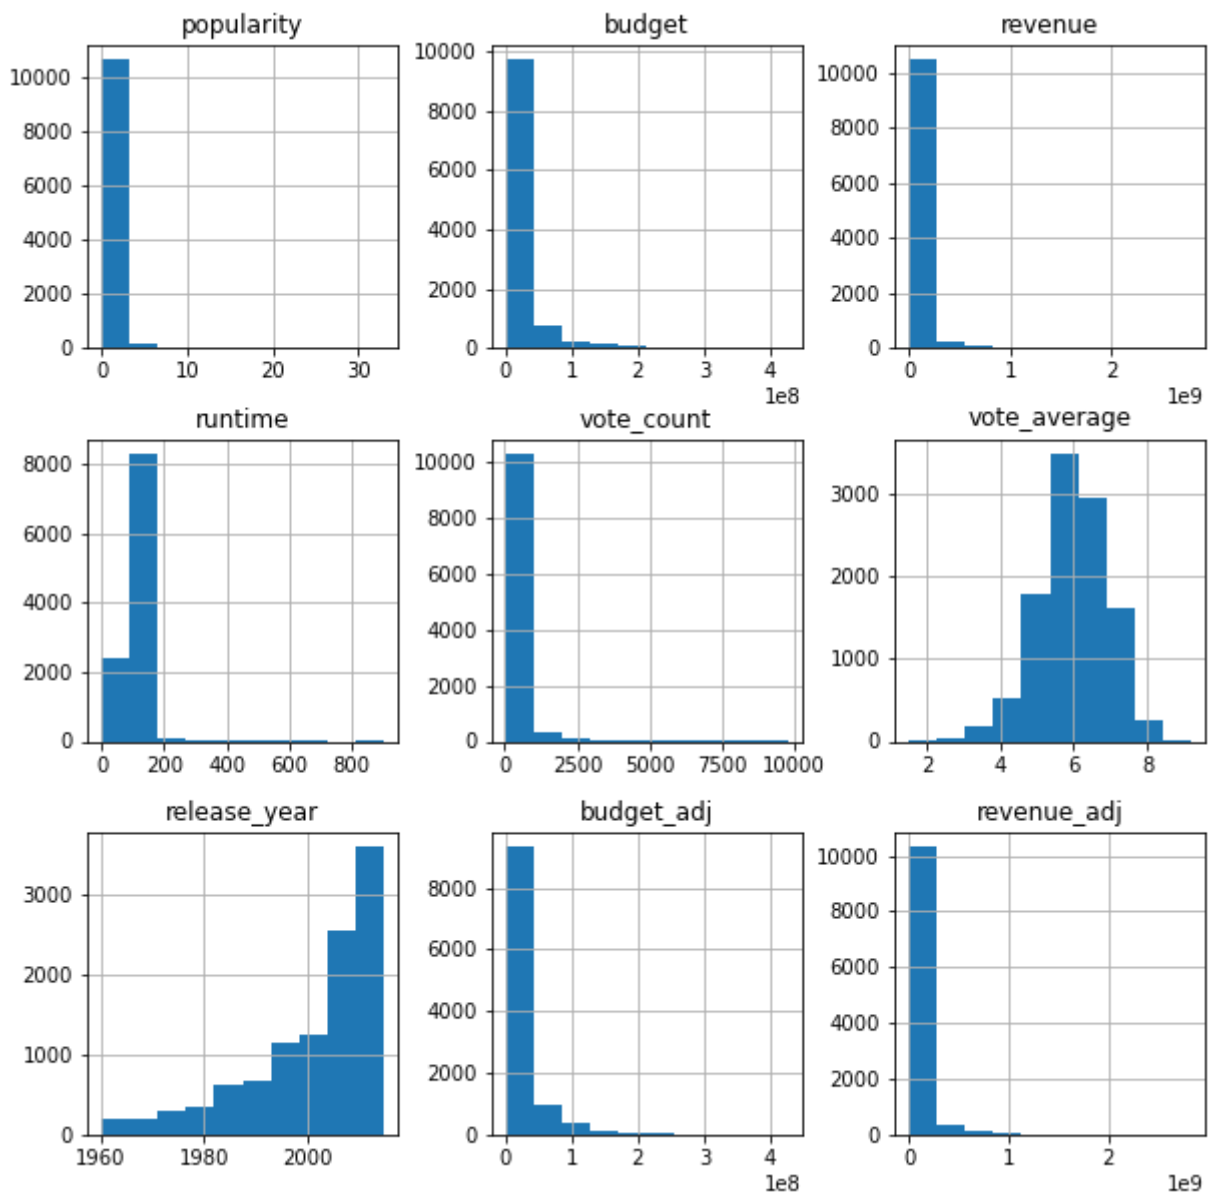
```
In [8]: # count zero values on `budget`, `revenue` and `popularity`
print('Zero values of Budget: ', df.query('budget == 0').count()['budget'])
print('Zero values of Revenue: ', df.query('revenue == 0').count()['revenue'])

Zero values of Budget:  5696
Zero values of Revenue: 6016
```

```
In [9]: # to check if popularity ranges mostly > 0 and less than 10
df.query('popularity < 10')['popularity']
```

```
Out[9]: 4          9.335014
5          9.110700
6          8.654359
7          7.667400
8          7.404165
...
10861      0.080598
10862      0.065543
10863      0.065141
10864      0.064317
10865      0.035919
Name: popularity, Length: 10855, dtype: float64
```

```
In [10]: df.iloc[:, 1:].hist(figsize=(10,10));
```



Data is described according to the above histograms. Which means, only visible value ranges is considered. We will check more exact values in the following.

- It seems from below that `popularity` table's almost all values ranges from 0 to 10.
- `Budget` table ranges from 0 million to 300 million. Currency name is not available, but should be dollar I think.
- `Revenue` table ranges from 0 billion to 1 billion. Currency name is not available, but should be dollar I think.
- `Runtime` ranges mostly from 0 to 200

- vote_average seems almost symmetric
- and so on ...

```
In [11]: # to determine, if release_year have missing value
null_data_release_year = df[df['release_year'].isnull()]
null_data_release_year.head(3)
```

Out[11]:

id	imdb_id	popularity	budget	revenue	original_title	cast	homepage	director	tagline	...	ove
----	---------	------------	--------	---------	----------------	------	----------	----------	---------	-----	-----

0 rows × 21 columns

```
In [12]: # to determine, if release_year have missing value
null_data_genres = df[df['genres'].isnull()]
# null_data_genres.shape
null_data_genres.head(2)
```

Out[12]:

id	imdb_id	popularity	budget	revenue	original_title	cast	homepage
424	363869	tt4835298	0.244648	0	0	Belli di papÃ	Diego Abatantuono Matilde Giolì Andrea Pisani ...
620	361043	tt5022680	0.129696	0	0	All Hallows' Eve 2	NaN

2 rows × 21 columns

Data Cleaning (Dropping, unique, merging, grouping, duplicate values)

Here, I decided to remove values of budget and revenue that equals to 0 at the same time. However, I will keep rows if one of them is not equals to zero

```
In [13]: # deleteing rows when budget and revenue eugals to zero
indexes_to_delete= df.query('budget == 0 and revenue == 0').index
df.drop(indexes_to_delete, inplace = True)
df.shape
```

Out[13]: (6165, 21)

Here, I will asses mean value of budget when it is equal to 0 ; However, I decided to give budget value to its revenue for avoiding big difference in values. I mean, I thought movie's profit at least should equal to its budget .)


```
In [14]: # Replacing 0 budget values with mean of all expenditure
df['budget'] = df['budget'].replace(0, df['budget'].mean())

# Assesing budget value to revenue when it is equal to zero
for index, row in df.iterrows():
    if df['revenue'][index] == 0:
        df['revenue'][index] = df['budget'][index]
```

```
In [15]: # Chech if zero values on `budget` and `revenue` exists
print('Zero values of Budget: ', df.query('budget == 0').count()['budget'])
print('Zero values of Revenue: ', df.query('revenue == 0').count()['revenue'])

Zero values of Budget: 0
Zero values of Revenue: 0
```

To answer Q1(Question - 1), I decided to transform given dataframe to `df_q1` which will be comfortable to draw barplot. The answer of the question 1 can be found in `release_year` and `genres`. Therefore, I did some cleaning process as following to above mentioned columns.

```
In [16]: # to select non-null genres with release_year for Question 1 (Q1)
df_q1 = df.dropna(subset=['genres'])[['release_year', 'genres']]

# To determine the non-null values lengh
print('Lenght of Q1 dataset : ', len(df_q1))
df_q1.head(4)
```

Lenght of Q1 dataset : 6164

Out[16]:

	release_year	genres
0	2015	Action Adventure Science Fiction Thriller
1	2015	Action Adventure Science Fiction Thriller
2	2015	Adventure Science Fiction Thriller
3	2015	Action Adventure Science Fiction Fantasy

```
In [17]: # Finding the unique values of the year
df_q1['release_year'].unique()
```

```
Out[17]: array([2015, 2014, 1977, 2009, 2010, 1999, 2001, 2008, 2011, 2002, 1994,
                2012, 2003, 1997, 2013, 1985, 2005, 2006, 2004, 1972, 1980, 2007,
                1979, 1984, 1983, 1995, 1992, 1981, 1996, 2000, 1982, 1998, 1989,
                1991, 1988, 1987, 1968, 1974, 1975, 1962, 1964, 1971, 1990, 1961,
                1960, 1976, 1993, 1967, 1963, 1986, 1973, 1970, 1965, 1969, 1978,
                1966])
```

```
In [18]: # Grouping all other genres in the same year by using `|` character
df_q1['genres'] = df_q1.groupby('release_year')['genres'].transform(lambda x:
df_q1.head()
```

Out[18]:

	release_year	genres
0	2015	Action Adventure Science Fiction Thriller Acti...
1	2015	Action Adventure Science Fiction Thriller Acti...
2	2015	Action Adventure Science Fiction Thriller Acti...
3	2015	Action Adventure Science Fiction Thriller Acti...
4	2015	Action Adventure Science Fiction Thriller Acti...

```
In [19]: #dropping duplicate values after merging dataset of `df_q1`
df_q1 = df_q1.drop_duplicates()

# sorting values to draw barplot in ascending order
df_q1 = df_q1.sort_values(by=['release_year'])
```

```
In [20]: # to check if release_year is sorted
df_q1.head(4)
```

Out[20]:

	release_year	genres
10141	1960	Drama Horror Thriller Action Adventure Western...
10110	1961	Adventure Animation Comedy Family Comedy Drama...
9849	1962	Adventure Action Thriller Adventure Drama Hist...
10438	1963	Action Thriller Adventure Horror Comedy Myster...

```
In [21]: # to check if genres is merged correctly by manually checking original data
df_q1['genres'][10141]
```

Out[21]: 'Drama|Horror|Thriller|Action|Adventure|Western|Action|Drama|History|Comedy|Drama|Romance|Thriller|Adventure|Fantasy|Science Fiction|Romance|Comedy|Drama|Romance|Horror|Thriller|Comedy|Horror|Science Fiction|Comedy|Family|Comedy|Romance'

```
In [22]: # shape should be equal to df_q1['release_year'].unique() value, manually c
df_q1.shape
```

Out[22]: (56, 2)

Exploratory Data Analysis

Research Question 1 (Which genres are most popular from year to year?)

Importance of the Research Question 1

By analyzing the question, we can explore which genres are in demand from year to year as well as overall. I am also going to determine top 2 genres every year to know more insights.

To explore, I am going to work with genres and release_year columns. To Explore Question 1 visually, I have to split genres by | character and should create columns to each genre with argument count values.

Mainly, I am planning to use barplot and piechart in my statistics result.

Function `get_genre_frequency` splits the genres and return it as `numpy.ndarray` like tuple. So, it returns `unique` - unique genre name and `counts` - counts the number of its appearance

```
In [23]: def get_genre_frequency(genres):
         array = np.array(genres.split('|'))
         (unique, counts) = np.unique(array, return_counts=True)
         return np.asarray((unique, counts)).T
```

Function `create_each_genre_columns`. Here, I am transferring all genres values into columns which is resulted from `get_genre_frequency` function.

```
In [24]: def create_each_genre_columns(np_genre_frequencies, df, ind):
         for unique, counts in np_genre_frequencies:
             if unique not in df.columns:
                 df[unique] = 0
             df[unique][ind] = counts
         return df
```

Function `create_the_most_popular_genre_columns` is created to visualize the winner of the genres from year to year.

```
In [25]: def create_the_most_popular_genre_columns(np_genre_frequencies, df, ind):
         df['popular_genre_count'] = df.iloc[:, 4:].max(axis=1)
         df['popular_genre_name'] = df.iloc[:, 4:].idxmax(axis=1)
         return df
```

Here I am running above 3 functions to create `df_q1` by considering index.

```
In [26]: df_q1['popular_genre_name'] = 'UNKNOWN'
df_q1['popular_genre_count'] = 0
for ind in df_q1.index:
    np_genre_frequencies = get_genre_frequency(df_q1['genres'][ind])
    df_q1 = create_each_genre_columns(np_genre_frequencies, df_q1, ind)
    df_q1 = create_the_most_popular_genre_columns(np_genre_frequencies, df_

df_q1.head(10)
```

Out[26]:

	release_year	genres	popular_genre_name	popular_
10141	1960	Drama Horror Thriller Action Adventure Western...	Comedy	
10110	1961	Adventure Animation Comedy Family Comedy Drama...	Drama	
9849	1962	Adventure Action Thriller Adventure Drama Hist...	Drama	
10438	1963	Action Thriller Adventure Horror Comedy Myster...	Adventure	
9881	1964	Adventure Action Thriller Drama Comedy War Com...	Drama	
10689	1965	Adventure Action Thriller Drama Family Music R...	Action	
10820	1966	Animation Family Comedy Drama Drama Family Adv...	Adventure	
10398	1967	Family Animation Adventure Comedy Drama Romanc...	Drama	
9719	1968	Science Fiction Mystery Adventure Adventure Sc...	Drama	
10724	1969	Adventure Action Thriller History Drama Wester...	Drama	

10 rows x 24 columns

```
In [27]: # Renaming release_year to Years for convinience
df_q1.rename(columns={'release_year': 'Years'}, inplace=True)
df_q1.head(3)
```

Out[27]:

	Years	genres	popular_genre_name	popular_genre_
10141	1960	Drama Horror Thriller Action Adventure Western...	Comedy	
10110	1961	Adventure Animation Comedy Family Comedy Drama...	Drama	
9849	1962	Adventure Action Thriller Adventure Drama Hist...	Drama	

3 rows x 24 columns

```
In [28]: # This cells code has been copied from here
# [https://towardsdatascience.com/reordering-pandas-dataframe-columns-thumbs]
# Aim of this function is to change the position of the columns
def movecol(df, cols_to_move=[], ref_col='', place='After'):

    cols = df.columns.tolist()
    if place == 'After':
        seg1 = cols[:list(cols).index(ref_col) + 1]
        seg2 = cols_to_move
    if place == 'Before':
        seg1 = cols[:list(cols).index(ref_col)]
        seg2 = cols_to_move + [ref_col]

    seg1 = [i for i in seg1 if i not in seg2]
    seg3 = [i for i in cols if i not in seg1 + seg2]

    return(df[seg1 + seg2 + seg3])
```

```
In [29]: # I am changing position of `Years` column with `popular_genre_count` column
df_q1 = movecol(df_q1,
                 cols_to_move=['popular_genre_count', 'Years'],
                 ref_col='popular_genre_name',
                 place='After')
df_q1.head(3)
```

Out[29]:

	genres	popular_genre_name	popular_genre_count
10141	Drama Horror Thriller Action Adventure Western...	Comedy	5
10110	Adventure Animation Comedy Family Comedy Drama...	Drama	9
9849	Adventure Action Thriller Adventure Drama Hist...	Drama	10

3 rows x 24 columns

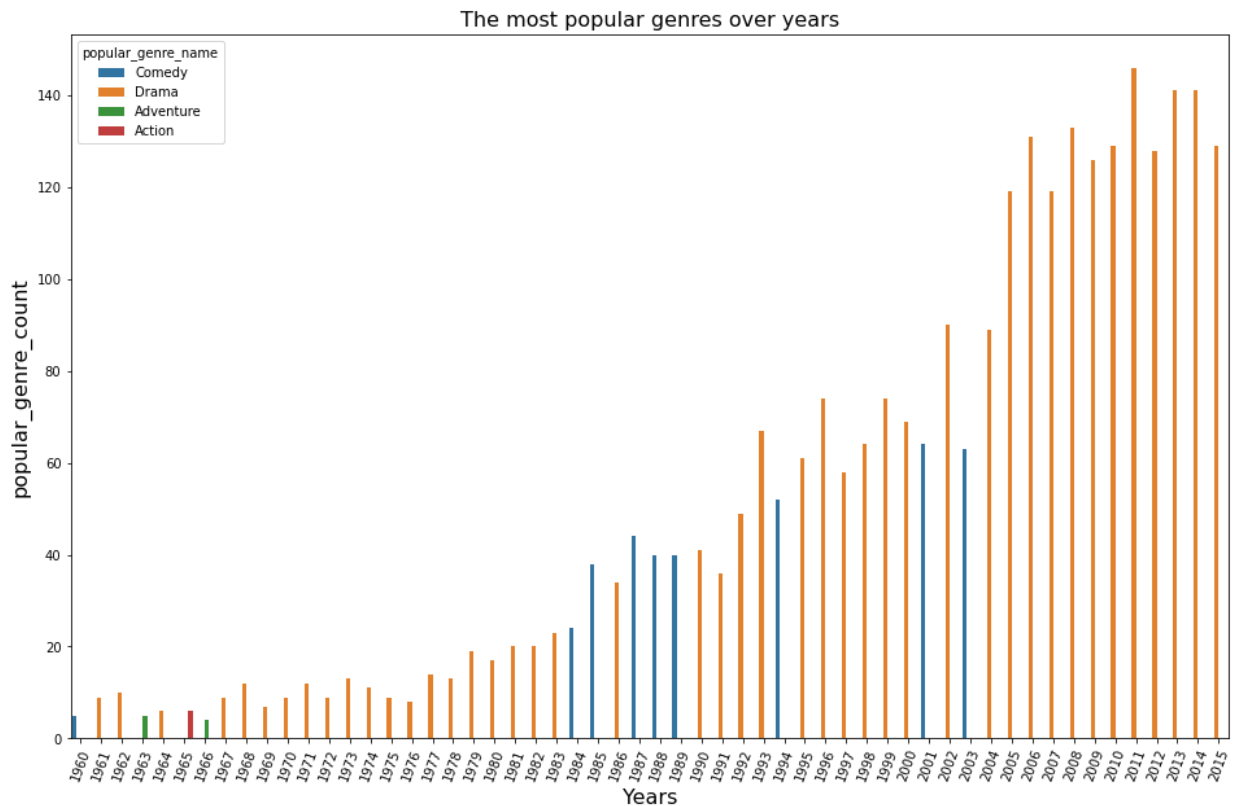
```
In [30]: # to count number of unique genre movies in the dataset, manually
df_q1.columns
```

```
Out[30]: Index(['genres', 'popular_genre_name', 'popular_genre_count', 'Years',
                'Action', 'Adventure', 'Comedy', 'Drama', 'Family', 'Fantasy',
                'History', 'Horror', 'Romance', 'Science Fiction', 'Thriller',
                'Western', 'Animation', 'Crime', 'Music', 'War', 'Foreign', 'Myste
ry',
                'Documentary', 'TV Movie'],
                dtype='object')
```

Answer to Q1 (Question 1).

It is obvious that Drama genres movies is the most popular from the below bar plot. However, rarely we can see comedy genre movies can beat the Drama genre movies.


```
In [31]: # Here I am using `Years`, `popular_genre_count` and `popular_genre_name` c
fig, ax1 = plt.subplots(figsize=(16, 10))
ax1.set_title('The most popular genres over years', fontsize=16)
ax1.set_xlabel('Years', fontsize=16)
ax1.set_ylabel('Count', fontsize=16)
ax1 = sns.barplot(x='Years', y='popular_genre_count', hue='popular_genre_na
plt.xticks(rotation=70)
plt.show()
sns.despine(fig)
```



According to the bar plot above Drama genre is the most popular genres over years, whereas you can rarely see some comedy, adventure and Action genre movies winning.

```
In [32]: # slicing df_q1 dataframe for pie chart data
df_q1_1 = df_q1.loc[:, 'Action': 'TV Movie']
df_q1_1 = df_q1_1.sum(axis = 0)
```

```
In [33]: #Using matplotlib's pie chart
pie, ax = plt.subplots(figsize=[10,6])
labels = list(df_q1_1.index.values)
explode = np.full(shape=len(labels), fill_value=0.1, dtype=np.float64)
plt.pie(x=df_q1_1, autopct="%.1f%%", explode=explode, labels=labels, pctdis
plt.title("All Genres share", fontsize=14);
pie.savefig("all_genres.png")
```



According to the pie chart above Drama genre is the most popular genres, the second is comedy , the third place is Thriller and follows Action , Adventure and so on.

To explore all other genres, I decided to draw bar plot by slicing the df_q1 dataframe within 15 years

```
In [34]: # to change the width of barplot
def change_width(ax, new_value) :
    for patch in ax.patches :
        current_width = patch.get_width()
        diff = current_width - new_value

        ax.annotate(format(patch.get_height(), '.1f'),
                    (patch.get_x() + patch.get_width() / 2., patch.get_
                     ha='center', va='center',
                     xytext=(0, 7), textcoords='offset points')

        # change the bar width
        patch.set_width(new_value)

        # I recenter the bar
        patch.set_x(patch.get_x() + abs(diff))
```

```
In [35]: # to draw the bar plot by grouping them yearly
def draw_bar_plot(title, df):
    fig, ax1 = plt.subplots(figsize=(14, 10))
    ax1.set_title(title, fontsize=16)
    ax1.set_xlabel('Years', fontsize=16)
    ax1.set_ylabel('Count', fontsize=16)
    tidy = df.melt(id_vars='Years').rename(columns=str.title)
    ax1 = sns.barplot(x='Years', y='Value', hue='Variable', data=tidy)
    plt.xticks(rotation=70)

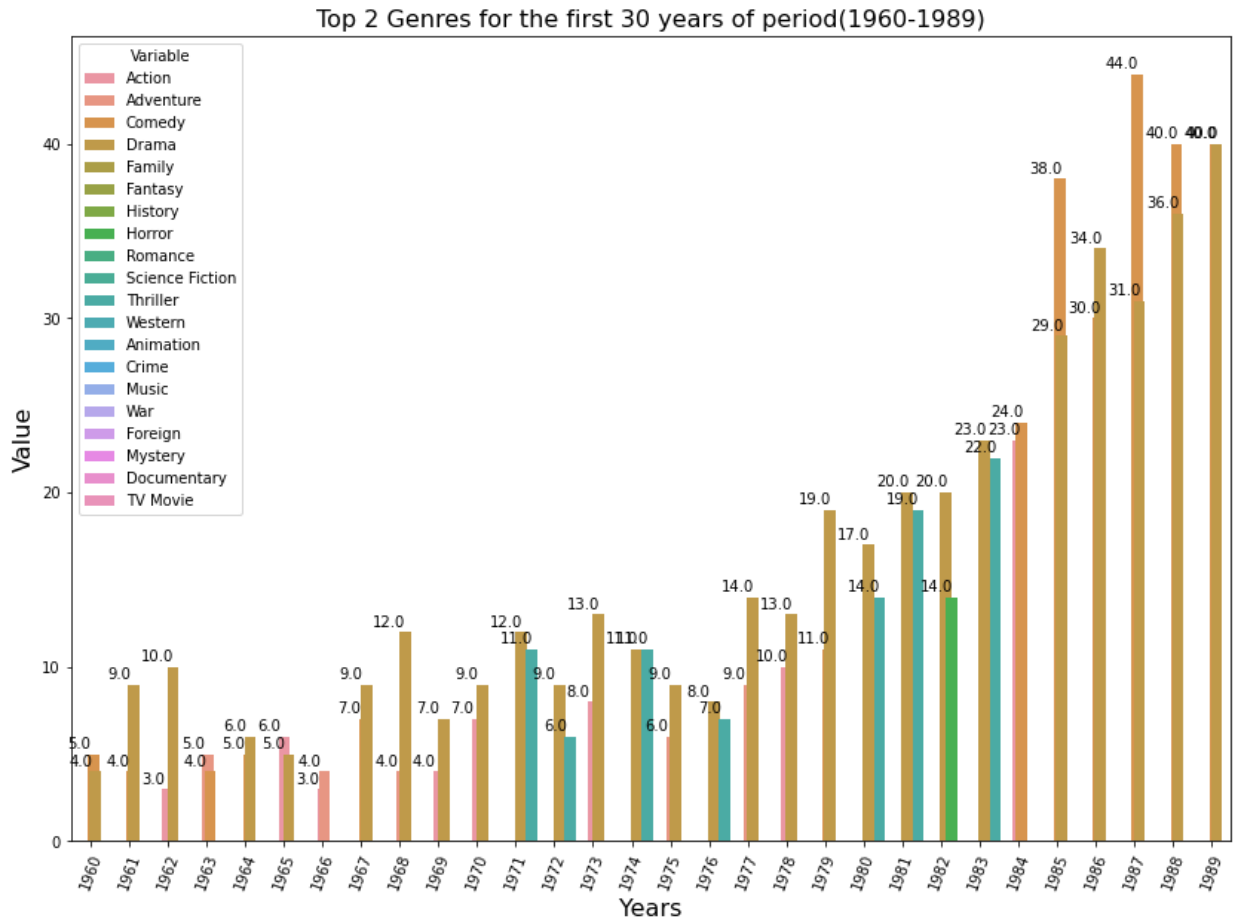
    change_width(ax1, .3)
    plt.show()
    sns.despine(fig)
```

```
In [36]: def get_top_three_genres(df):
    for index, row in df.iloc[:, 1:].iterrows():
        n_largest = row.nlargest(2)
        genres_list = list(n_largest.index.values)
        genres_list.append('Years')
        for column in df:
            if (column not in genres_list):
                df[column][index] = None
    return df
```




```
In [37]: # Bar Plot of all Genres from 1960 to 2015
df_q1_1 = get_top_three_genres(df_q1.iloc[:30, 3:])
draw_bar_plot('Top 2 Genres for the first 30 years of period(1960-1989)', d

df_q1_2 = get_top_three_genres(df_q1.iloc[30:56,3:])
draw_bar_plot('Top 2 Genres for the second 24 years of period(1990-2015)',
```



Top 2 Genres for the second 24 years of period(1990-2015)



Above two barplot shows top 2 number of genre from year to year. number of movies created from year to year in all periods have increased considerably. Especially, Drama is increased considerably higher than other genres

Research Question 2 (Q2.What kinds of properties are associated with movies that have high revenues?)

Importance of the Research Question 2

In this research question, I am going to find out which properties are most important to make big profit. For example, does investing high amount of budget will be reason to make high revenue? or Does length(runtime) of movie realted? and etc...

- In the analysis process, firstly, I will find out correlation between continuous number columns. Regression plot of seaborn will be appropriate;
- Then I will explore it using Seaborn's heatmap plot

To answer the question following posts' idea is partially used

<https://towardsdatascience.com/correlation-is-simple-with-seaborn-and-pandas-28c28e92701e>
<https://towardsdatascience.com/correlation-is-simple-with-seaborn-and-pandas-28c28e92701e>.

None of code is copied!

```
In [38]: # to see the highest revenue value
df_q2 = df.sort_values(by='revenue', ascending=False)
df_q2.head(2)
```

Out[38]:

	id	imdb_id	popularity	budget	revenue	original_title	cast
1386	19995	tt0499549	9.432768	237000000.0	2781505847	Avatar	Sam Worthington Zoe Saldana Sigourney Weaver S...
3	140607	tt2488496	11.173104	200000000.0	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...

2 rows x 21 columns



```
In [39]: def plot_correlation(xPlot, YPlot, df, title, isCurrency):
fig, ax1 = plt.subplots(figsize=(7, 4))
ax1 = sns.regplot(x=xPlot, y=YPlot, data=df);
ax1.set_title(title, fontsize=16)
ax1.set(xlim = (min(df[xPlot]),max(df[xPlot])))
ax1.set(ylim = (min(df[YPlot]),max(df[YPlot])))



xlabels = ['${:,.2f}'.format(x) + 'B' for x in ax1.get_xticks()/1000000]
ticks_loc = ax1.get_xticks().tolist()
ax1.xaxis.set_major_locator(ticker.FixedLocator(ticks_loc))
ax1.set_xticklabels(xlabels)
ax1.xaxis.set_tick_params(which='major', labelcolor='green')

if isCurrency:
    ylabels = ['${:,.2f}'.format(x) + 'B' for x in ax1.get_yticks()/100]
    ticks_loc = ax1.get_yticks().tolist()
    ax1.yaxis.set_major_locator(ticker.FixedLocator(ticks_loc))
    ax1.set_yticklabels(ylabels)
    ax1.yaxis.set_tick_params(which='major', labelcolor='green',
                              labelleft=True, labelright=False)

plt.show()
```

```
In [40]: plot_correlation('revenue', 'budget', df_q2, 'Revenue vs Budget in Billions',  
plot_correlation('revenue', 'popularity', df_q2, 'Revenue vs Popularity', F  
plot_correlation('revenue', 'vote_count', df_q2, 'Revenue vs Vote Count', F  
plot_correlation('revenue', 'runtime', df_q2, 'Revenue vs Runtime', False);
```





- It seems Revenue vs Budget and Revenue vs Vote Count have positive correlation whereas others not much correlated.
- It seems number of movies that spent more than 200 million dollar budget is very scarce.
- It seems most of the movies doesn't earn more than 4000 vote_count .
- It seems Runtime doesn't much related to movie's revenue

Let's check the following cells with exact numbers.

```
In [41]: # to calculate correlation between continuous numbers
correlations = df_q2.corr()
```

```
In [42]: correlations.iloc[1:, 1:]
```


```
Out[42]:
```

	popularity	budget	revenue	runtime	vote_count	vote_average	release_year	bu
popularity	1.000000	0.478008	0.641816	0.188585	0.794867	0.286433	0.145786	
budget	0.478008	1.000000	0.702172	0.242477	0.578710	0.084338	0.197265	
revenue	0.641816	0.702172	1.000000	0.222407	0.773191	0.215797	0.101229	
runtime	0.188585	0.242477	0.222407	1.000000	0.227404	0.304330	-0.102140	
vote_count	0.794867	0.578710	0.773191	0.227404	1.000000	0.335786	0.166422	
vote_average	0.286433	0.084338	0.215797	0.304330	0.335786	1.000000	-0.131538	
release_year	0.145786	0.197265	0.101229	-0.102140	0.166422	-0.131538	1.000000	
budget_adj	0.453438	0.920227	0.679727	0.301063	0.540527	0.092213	0.048262	
revenue_adj	0.580560	0.559666	0.908316	0.238660	0.682493	0.252553	-0.085461	

Let's look at the `revenue` column from above and discuss:

- It is the most probable that `spending` high budget resulted high `revenue`.
- The second most probable association is `vote_count`. So, if `vote_count` is high, then movie can earn high amount of money.
- The third most probable association is `popularity` for `revenue`. So, It can be considered that if the popularit value is high then they earned high revenue
- `runtime` and `vote_average` doesn't seem to be related to the `revenue`

```
In [43]: # Code idea is copied from [https://medium.com/@szabo.bibor/how-to-create-a-
plt.figure(figsize=(16, 6))
heatmap = sns.heatmap(correlations.iloc[1:, 1:], vmin=-1, vmax=1, annot=True)
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':12}, pad=12);
plt.show()
```



Research Question 3 (Q3. Which movies made the most profit, yearly?)

Importance of the Research Question 3

In this research question, I am going to find out movies which made the highest amount of profit from year to year? By doing yearly, I will be able to know top profitable movie types.

```
In [44]: df_q3 = pd.DataFrame()
df_q3['profit'] = df['revenue'] - df['budget']
df_q3['year'] = df['release_year']
df_q3['movie_name'] = df['original_title']
df_q3.head()
```

Out[44]:

	profit	year	movie_name
0	1.363529e+09	2015	Jurassic World
1	2.284364e+08	2015	Mad Max: Fury Road
2	1.852382e+08	2015	Insurgent
3	1.868178e+09	2015	Star Wars: The Force Awakens
4	1.316249e+09	2015	Furious 7

```
In [45]: # df_q3 = df_q3.groupby(['year'], sort=True)['profit'].max()
idxs = df_q3.groupby(['year'], sort=False)['profit'].transform(max) == df_q3['profit']
df_q3 = df_q3[idxs]
df_q3 = df_q3.sort_values(by='year', ascending=True, na_position='first')
df_q3 = df_q3.reset_index()
df_q3.head(5)
```

Out[45]:


	index	profit	year	movie_name
0	10143	48000000.0	1960	Spartacus
1	10110	211880014.0	1961	One Hundred and One Dalmatians
2	9849	58500000.0	1962	Dr. No
3	10438	76398765.0	1963	From Russia With Love
4	9881	121400000.0	1964	Goldfinger

```
In [46]: fig, ax1 = plt.subplots(figsize=(16, 10))
ax1.set_title('Movies` profit', fontsize=16)
ax1 = sns.barplot(x='year', y='profit', data=df_q3)
plt.xticks(rotation=70)

ylab = ['${:,.2f}'.format(x) + 'B' for x in ax1.get_yticks()/1000000000]
ticks_loc = ax1.get_yticks().tolist()
ax1.yaxis.set_major_locator(ticker.FixedLocator(ticks_loc))
ax1.set_yticklabels(ylab)
ax1.yaxis.set_tick_params(which='major', labelcolor='green',
                           labelleft=True, labelright=False)

def autolabel(rects):
    for i in range(0, len(rects)):
        height = rects[i].get_height()
        ax1.text(rects[i].get_x() + rects[i].get_width() / 2.,
                  1.01 * height,
                  df_q3.iloc[i]['movie_name'],
                  ha='center', va='bottom', rotation=90, color='black', fontd

autolabel(ax1.patches)
plt.show()
```



Analysis Shortcoming & Data Limitations

Please note that the report is expected to include a discussion about the limitations and shortcomings of the analysis and the dataset. For example, Do you think that the number of samples is a good representative of the population? Do you think that missing values might influence the results of the analysis? The analysis includes comparisons between different categories, you can discuss the importance of statistical test to appreciate the significance of the differences. Optionally you can even implement the appropriate statistical test.

Size of the table: (10866 - rows, 21 - columns)

Representatation of samples can be sometimes challenging to prcoess, for example, splitting genres by | character and transforming each unique value(with its count) to column dataframe can increase time complexity.

More than half of the data for column budget and revenue equal to zero, which requires dropping and replacing. For instance, I dropped given the dataframe rows when both of the columns equal to zero . However, if one of them is not equal to zero, I took mean for budget and for revenue equal to zero, I transformed budget value.

I believe, there are many aspects to explore the dataset. However, my solution above should work pretty well for genres, revenue and movies name.

Conclusions

Following summarizations I get from three research questions above

- Drama genre movies is the most popular from the below bar plot. However, rarely we can see comedy genre movies can beat the Drama genre movies.
- Action , Thriller , Romance genres are the next most made movies after Drama .
- Overall, there 20 unique genre movies in the dataset
- Correaltions beetwen revenue and other tables are as following:
 - Correlation beetwen revenue and vote_count is the highest val = (0.79) ;
 - The second and third highest corresponds to budget(0.73) and popularity``(0.66) , respectively;
 - It seems from the given dataset, runtime(0.16) and vote_average(0.17) properties doesn't associated with revenue.
- Avatar moveie made the most profit, followed by Star Wars: The Force Awakens and Titanic

August 1, 01:29, 2021. Made by Sanatbek Matlatipov

